



PROJECT:-

HOUSING PROJECT

AUTHOR:-

ATUL DAHIMAN

PROBLEM STATEMENT:-

Houses are one of the necessary need of each and every person around the globe and therefore housing and real estate market is one of the markets which is one of the major contributors in the world's economy. It is a very large market and there are various companies working in the domain. Data science comes as a very important tool to solve problems in the domain to help the companies increase their overall revenue, profits, improving their marketing strategies and focusing on changing trends in house sales and purchases. Predictive modelling, Market mix modelling, recommendation systems are some of the machine learning techniques used for achieving the business goals for housing companies. Our problem is related to one such housing company. A US-based housing company named Surprise Housing has decided to enter the Australian market. The company uses data analytics to purchase houses at a price below their actual values and flip them at a higher price. For the same purpose, the company has collected a data set from the sale of houses in Australia. The data is provided in the CSV file below.

The company is looking at prospective properties to buy houses to enter the market. You are required to build a model using Machine Learning in order to predict the actual value of the prospective properties and decide whether to invest in them or not. For this company wants to know:

- Which variables are important to predict the price of variable?
- How do these variables describe the price of the house?

Business Goal:

You are required to model the price of houses with the available independent variables. This model will then be used by the management to understand how exactly the prices vary with the variables. They can accordingly manipulate the strategy of the firm and concentrate on areas that will yield high returns. Further, the model will be a good way for the management to understand the pricing dynamics of a new market.

IMPORTING LIBRARIES AND LOAD DATA:-

import libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings("ignore")
```

load data

```
In [2]: train_data=pd.read_csv(r"C:\Users\ankus\Downloads\train.csv")
test_data=pd.read_csv(r"C:\Users\ankus\Downloads\test.csv")
train_data.head()
```

```
Out[2]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	Mo
0	127	120	RL	NaN	4928	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1	889	20	RL	95.0	15865	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
2	793	60	RL	92.0	9920	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
3	110	20	RL	105.0	11751	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0	
4	422	20	RL	NaN	16635	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	

5 rows × 81 columns

There are two data set one is train and other is test dataset . We have to make prediction on the test dataset using train dataset

Few libraries are used initially for different purpose like:-

- Numpy:- for mathematical operation
- Pandas:- for data preprocessing
- Matplotlib and Seaborn:- for data visualization
- Warnings:- for some warnings

As most of the columns are non numerical so we will do label encoding first

LABEL ENCODING:-

Label encoding

```
In [6]: #for train data
from sklearn.preprocessing import LabelEncoder

train_data = train_data.copy()
label_encoder = LabelEncoder()
for column in train_data.columns:
    train_data[column] = label_encoder.fit_transform(train_data[column])
train_data.head()
```

```
Out[6]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Cond
0	94	11	3	106	80	1	2	0	3	0	4	0	13	2	
1	720	0	3	65	808	1	2	0	3	0	4	1	12	2	
2	642	5	3	62	449	1	2	0	3	0	1	0	15	2	
3	79	0	3	75	632	1	2	0	3	0	4	0	14	2	
4	341	0	3	106	821	1	2	0	3	0	2	0	14	2	

Now we can clearly see that after doing label encoding all the columns have changed into numerical column. Same as we have to do with test data also.

CHECK FOR NULL VALUES AND DATA DESCRIPTION:-

```
In [11]: train_data.isnull().sum().sum()
Out[11]: 0

In [12]: test_data.isnull().sum().sum()
Out[12]: 0

In [11]: #check data description
train_data.describe()

Out[11]:
   Id  MSSubClass  MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  LandContour  Utilities  LotConfig  LandSlope
  count  1168.000000  1168.000000  1168.000000  1168.000000  1168.000000  1168.000000  1168.000000  1168.000000  1168.0  1168.000000  1168.000000
  mean  583.500000  4.166096  3.013699  52.454623  414.643836  0.996575  1.898973  1.938356  2.773973  0.0  3.004281  0.064212
  std  337.316864  4.139986  0.633120  31.378301  249.993254  0.058445  0.401453  1.412262  0.710027  0.0  1.642667  0.284088
  min  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.0  0.000000  0.000000
  25%  291.750000  0.000000  3.000000  30.000000  197.750000  1.000000  2.000000  0.000000  3.000000  0.0  2.000000  0.000000
  50%  583.500000  4.000000  3.000000  45.000000  407.500000  1.000000  2.000000  3.000000  3.000000  0.0  4.000000  0.000000
  75%  875.250000  6.000000  3.000000  70.000000  618.250000  1.000000  2.000000  3.000000  3.000000  0.0  4.000000  0.000000
  max  1167.000000  14.000000  4.000000  106.000000  891.000000  1.000000  2.000000  3.000000  3.000000  0.0  4.000000  2.000000
```

data.isnull().sum() shows that there is no null values in the data .

While data description shows some calculation about the data like:-mean,median,mode,std and some other calculation

On the top of the data description we can see that all the columns have same value which also defines that there is no null value

SEPERATING LABELS AND FEATURES:-

seperating label and features

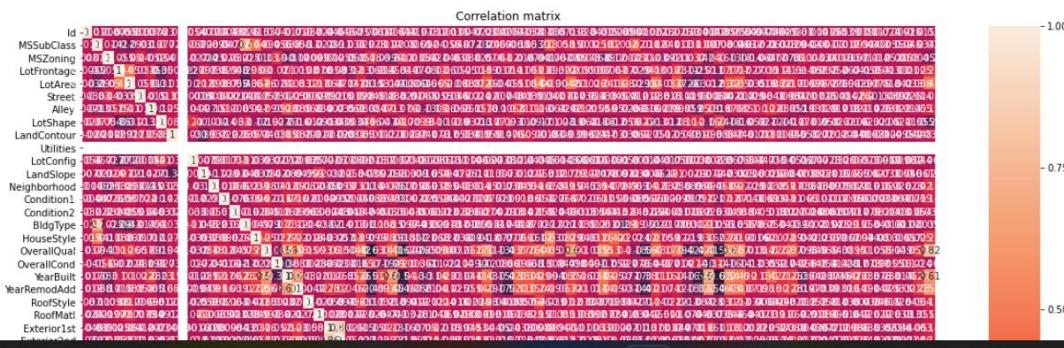
```
In [12]: x=train_data.drop(['SalePrice'],axis=1)
y=train_data['SalePrice']
```

Above we have separate the data set into label and features. As we have to predict the sales price of houses so sales price is the target variable.

CORRELATION:-

```
In [16]: #correlation using heatmap
plt.figure(figsize=[20,20])

#plot the correlation matrix
sns.heatmap(train_data.corr(), annot=True)
plt.title("Correlation matrix")
plt.show()
```



Above diagram shows the correlation among all the columns.

As here are so much column so we cannot easily identify the correlation among all the columns so lets print the correlation table.

```
In [15]: train_data.corr()
```

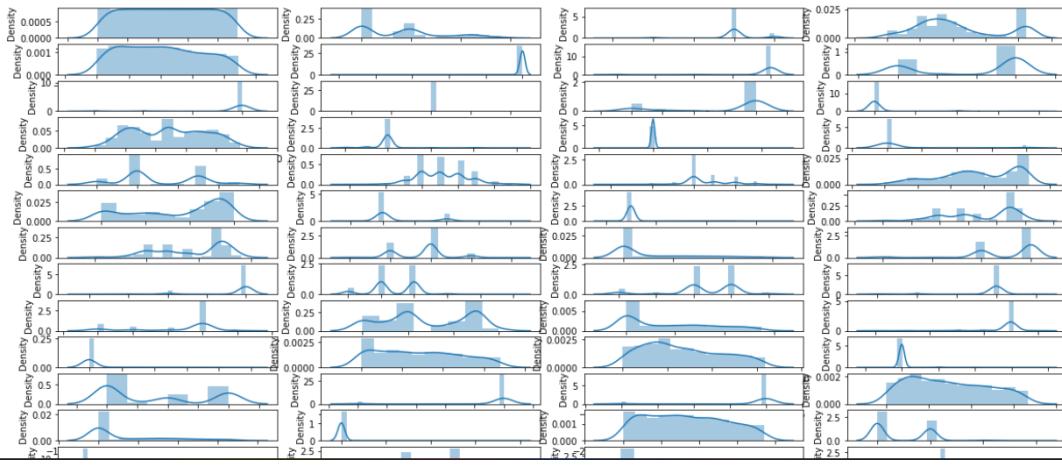
	BsmthHalfBath	0.024663	0.006345	0.000975	0.026097	0.013788	0.005111	-0.017926	0.013725	NaN	-0.016040	0.074245
FullBath	-0.015459	0.178100	-0.188837	0.103484	0.236095	0.033208	0.059465	-0.175359	0.047016	NaN	0.000177	-0.077542
HalfBath	-0.028426	0.197825	-0.121209	0.066610	0.130289	0.045146	0.050163	-0.122586	0.030155	NaN	-0.020604	-0.000910
BedroomAbvGr	0.009781	0.031998	-0.000228	0.165149	0.277605	0.027690	-0.028218	-0.055556	-0.038239	NaN	-0.058060	-0.071538
KitchenAbvGr	0.001028	0.324594	0.026744	-0.029430	-0.028955	0.012304	-0.065582	0.082102	-0.067071	NaN	-0.002959	-0.033515
KitchenQual	-0.000217	-0.018333	0.111689	-0.041490	-0.152446	-0.012056	-0.015901	0.122861	0.027230	NaN	-0.013554	-0.006002
TotRmsAbvGrd	-0.001837	0.116792	-0.030274	0.182701	0.397428	0.038317	-0.002743	-0.126699	-0.050646	NaN	-0.043083	-0.061747
Functional	0.020059	0.040881	-0.091684	0.018003	-0.032568	-0.015309	0.018574	-0.012018	0.019517	NaN	-0.016756	-0.121209
Fireplaces	-0.024571	-0.010739	0.010658	0.231717	0.366189	0.010574	0.097905	-0.183316	-0.053860	NaN	-0.056197	0.100623
FireplaceQu	0.013160	0.031013	-0.003746	-0.106440	-0.263034	-0.018641	-0.059485	0.107408	0.038313	NaN	0.011536	-0.002928
GarageType	0.004098	0.108879	0.135666	-0.254413	-0.304995	-0.030325	-0.253752	0.196219	-0.102398	NaN	0.037598	0.019580
GarageYrBlt	-0.004159	0.097167	-0.239603	-0.011122	-0.011836	-0.000639	0.177902	-0.156218	0.093121	NaN	0.000228	-0.071504

Using above table we can easily identify the correlation among all the columns.

DATA DISTRIBUTION:-

```
In [17]: #Let's see how data is distributed for every column
plt.figure(figsize=(20,20))
plotnumber=1

for column in train_data:
    if plotnumber<=81:
        ax=plt.subplot(27,4,plotnumber)
        sns.distplot(train_data[column])
        plt.xlabel(column,fontsize=20)
    plotnumber+=1
plt.show()
```



Above code shows the data distribution among all the columns , the data distribution is looking not so good , let go ahead with checking skewness in the data distribution.

SKEWNESS:-

```
In [18]: #checking skewness
x.skew().sort_values()

Out[18]: Street          -17.021969
PoolQC         -15.983184
MiscFeature     -5.406583
Alley           -4.056922
Functional      -3.999663
GarageCond      -3.789819
SaleType         -3.660513
GarageQual      -3.505364
CentralAir       -3.475188
BsmtFinType2    -3.388419
PavedDrive      -3.274035
Landcontour     -3.125982
Electrical       -3.104209
BsmtCond        -2.927336
SaleCondition    -2.671829
ExterCond        -2.516219
Fence            -1.955758
ExterQual        -1.810843
MSZoning         -1.796785
KitchenQual      -1.408106
LotConfig         -1.118821
BsmtQual         -1.107099
BsmtExposure     -1.075098
FireplaceQu      -0.794843
Exteriorist      -0.612816
LotShape          -0.603775
GarageYrBlt      -0.602871
Exterior2nd      -0.592349
YearRemodAdd     -0.495864
YearBuilt         -0.448970
GarageCars        -0.358556
GarageFinish      -0.129987
BsmtFinType1     -0.019567
Foundation        -0.002761
```

As we can see that lots of column are highly skewed so we will try some method for handle the skewness of the data

```
In [19]: #handle the skewness of teh data using sqrt method
x=np.sqrt(x)
x.skew()
```

Here we use the np.sqrt() method for handling the skewness of the data.

DATA SCALING:-

Data Scaling

```
In [ ]: #data scaling
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
x_scaled=scaler.fit_transform(x)
```

```
In [ ]: x_scaled
```

vif score

```
In [ ]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif=pd.DataFrame()
vif[["vif"]]=variance_inflation_factor(x_scaled, i) for i in range(x_scaled.shape[1])]

vif[["Features"]]=x.columns

#let's check the values
vif
```

Scaled the features variables and then check the vif score for scaled data to check the multicollinearity using “variance inflation factor”

TRAINING PROCESS:-

Here we will use Linear Regression model for model training and testing

Start the training process and some important module for model training and then split the data in training and testing data set

Training Process

```
In [ ]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
from sklearn.metrics import r2_score,confusion_matrix,classification_report
from sklearn.model_selection import train_test_split

In [ ]: x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.25,random_state=6)
```

MODEL INSTANTIATING AND TRAINING

```
In [ ]: lr.fit(x_train,y_train)
In [ ]: #Adjusted r2 score
lr.score(x_train,y_train)*100
In [ ]: #let's check how well model fits the test data
lr.score(x_test,y_test)*100
```

let's plot and visualize

```
In [ ]: y_pred=lr.predict(x_test)
In [ ]: y_pred
In [ ]: plt.scatter(y_test,y_pred)
plt.xlabel('Actual Happiness Score')
plt.ylabel('Predicted Happiness Score')
plt.title('Actual vs Predicted Score')
plt.show()
```

Check r2 score for training and testing data and then plot the graph for actual and predicted result with scatter plot

After this do cross validation and hyper parameter tunning for the model

The screenshot shows a Jupyter Notebook interface with the title "WORLD HAPPINESS REPORT". The notebook contains two sections:

- Cross Validation:**

```
In [ ]: train_accuracy=lr.score(x_train,y_train)
test_accuracy=lr.score(x_test,y_test)
from sklearn.model_selection import cross_val_score,GridSearchCV
cv_score=cross_val_score(lr,x_scaled,y, cv=5)
cv_mean=cv_score.mean()
cv_mean
```
- HYPER PARAMETER TUNNING:**

```
In [ ]: from sklearn.ensemble import RandomForestRegressor
parameters=[{'criterion':['mse','mae'],'max_features':['auto','sqrt','log2']}
rf=RandomForestRegressor()
clf=GridSearchCV(rf,parameters)
clf.fit(x_train,y_train)
print(clf.best_params_)

In [ ]: rf=RandomForestRegressor(criterion="mse",max_features="auto")
rf.fit(x_train,y_train)
rf.score(x_train,y_train)
pred_decision=rf.predict(x_test)

rfs=r2_score(y_test,pred_decision)
print("R2 score : ",rfs*100)

rfscore=cross_val_score(rf,x_scaled,y, cv=5)
rfc=rfscore.mean()
print('cross_val_score : ',rfc*100)
```

Here we did the cross validation of the model and after do the hyper parameter tunning of the model.

We get the final accuracy of the model as 89%.

SAVE MODEL:-

save model for later use

```
In [49]: import pickle  
filename='Housing Project.pkl'  
pickle.dump(rf,open(filename,'wb'))
```

```
In [ ]:
```

prediction using saved model

```
In [50]: loaded_model=pickle.load(open('Housing Project.pkl','rb'))  
result=loaded_model.score(x_test,y_test)  
print(result*100)
```

```
89.73449826978447
```

```
In [51]: conclusion=pd.DataFrame([loaded_model.predict(test_data)[:,],pred_decision[:,]],index=["Original","Predicted"])  
conclusion
```

```
Out[51]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
Original	429.62	445.46	438.50	437.65	450.77	442.97	433.50	430.12	449.01	434.50	372.47	422.20	456.05	429.04	445.80	418.77	429.28	441.71	430.88
Predicted	115.53	183.75	307.27	435.77	174.40	462.76	222.68	238.57	337.64	262.92	399.04	392.98	197.02	219.54	443.28	344.81	189.92	373.09	298.36

```
In [ ]:
```

Save the model for late use to prediction.

SUMMARY:-

We started with the data exploration where we got a feeling for the dataset, checked about missing data and learned which features are important. During this process we used seaborn and matplotlib to do the visualizations. During the data pre-processing part.

Afterwards we started training model using Linear Regression and applied cross validation on it. Then we do the Hyper Parameter tunning of the model and get the final accuracy score for the model.

Later we save the model for later use for prediction.